

# Mooncake.jl

Will Tebbutt

2025-10-03

# Introduction

- ▶ Algorithmic Differentiation in Julia
- ▶ Who: Me (ATI)
- ▶ Who: Hong Ge (University of Cambridge, ATI)
- ▶ Who: Guillaume Dalle, Xianda Sun, and many others!

# Introduction

```
import DifferentiationInterface as DI
import Mooncake

backend = DI.AutoMooncake()
p = DI.prepare_gradient(f, backend, x)
DI.value_and_gradient(f, p, backend, x)
```

# Introduction

```
function f(x::AbstractArray{<:Real})  
    s = zero(eltype(x))  
    for n in eachindex(x)  
        s += cos(exp(x[n]))  
    end  
    return s  
end  
  
x = randn(1024)
```

Tool	Ratio	Language
Enzyme.jl	2.89	LLVM
Mooncake.jl	2.15	Julia
ReverseDiff.jl	7.41	Julia
Zygote.jl	427	Julia

Your numbers *will* differ.

# Introduction

```
function f(x::AbstractArray{<:Real})  
    s = zero(eltype(x))  
    for n in eachindex(x)  
        s += 5 * x[n]  
    end  
    return s  
end  
  
x = randn(1024)
```

Tool	Ratio	Language
Enzyme.jl	6.36	LLVM
Mooncake.jl	9.99	Julia
ReverseDiff.jl	73.0	Julia
Zygote.jl	5370	Julia

Your numbers *will* differ.

# Why Bother?

# Questions

Could we improve on `ReverseDiff` / `Zygote` while staying within the language?

- ▶ Testable?
- ▶ Composition “guarantees”?
- ▶ Good primal performance  $\implies$  good AD performance?
- ▶ Wider language support inc. mutation?

# AD Recap



# What?

- ▶  $f : \mathbb{R}^P \rightarrow \mathbb{R}^Q$ , differentiable at  $x$
- ▶  $\mathbf{J} \in \mathbb{R}^{Q \times P}$ ,  $\mathbf{J}_{q,p} := \partial f_q / \partial x_p$
- ▶ Forward Mode:  $\dot{x} \mapsto \mathbf{J}\dot{x}$ ,  $\dot{x} \in \mathbb{R}^P$
- ▶ Reverse Mode:  $\bar{y} \mapsto \mathbf{J}^\top \bar{y}$ ,  $\bar{y} \in \mathbb{R}^Q$

# How?

- ▶ Repeated application of the chain rule
- ▶  $f := f_2 \circ f_1 \implies \mathbf{J} = \mathbf{J}_2 \mathbf{J}_1$
- ▶ Forward Mode:  $(\dot{x}_2 \mapsto \mathbf{J}_2 \dot{x}_2) \circ (\dot{x}_1 \mapsto \mathbf{J}_1 \dot{x}_1)$
- ▶ Reverse Mode (Forward-Pass): construct + store  $\bar{y}_n \mapsto \mathbf{J}^\top \bar{y}_n$
- ▶ Reverse Mode (Reverse-Pass): load and apply:  
 $(\bar{y}_2 \mapsto \mathbf{J}_2 \bar{y}_2) \circ (\bar{y}_1 \mapsto \mathbf{J}_1 \bar{y}_1)$

# Testability and Tangents

# Testability and Tangents

- ▶  $f : \mathbb{R}^P \rightarrow \mathbb{R}^Q$ , differentiable
- ▶  $\mathbf{J}$ : Jacobian of  $f$  at some  $x$
- ▶  $\langle \bar{y}, \mathbf{J}\dot{x} \rangle = \langle \mathbf{J}^\top \bar{y}, \dot{x} \rangle$ ,  $\dot{x} \in \mathbb{R}^P$ ,  $\bar{y} \in \mathbb{R}^Q$
- ▶  $x$  isa  $P$ , then  $dx$  isa `tangent_type(P)`
- ▶ Ensure  $\langle \cdot, \cdot \rangle$  works for `tangent_type(P)` for all  $P$
- ▶ Define function: `test_rule(rng, f, x...)`

# Composition and Tangents

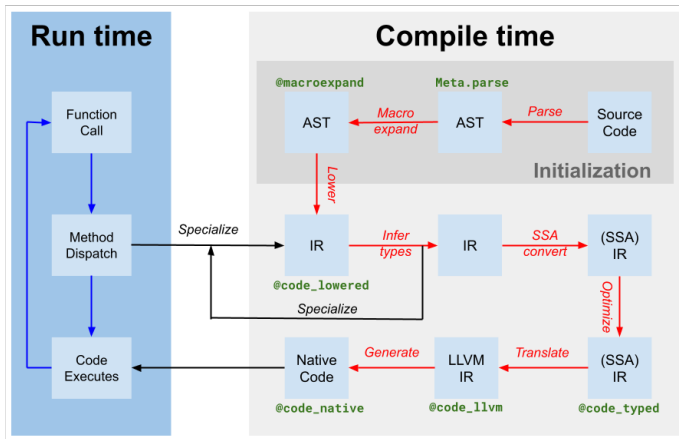
$$f = f_2 \circ f_1 \implies \mathbf{J}\dot{x} = \mathbf{J}_2(\mathbf{J}_1\dot{x})$$

```
function f(x)
    t = f_1(x)
    return f_2(t)
end
function rule(::typeof(f), (x, dx))
    (t, dt) = rule(f_1, (x, dx))
    return rule(f_2, (t, dt))
end
```

- ▶ Not formalised in Zygote
- ▶ Doable via `typeof(dt) = tangent_type(typeof(t))`

# Performance

# Transform Julia IR



<https://docs.julialang.org/en/v1/devdocs/jit/>



# Transform Julia IR

```
f(x) = sin(cos(x))  
Base.code_ircode_by_type(  
    Tuple{typeof(f),Float64}  
)[1][1]
```

```
julia> Base.code_ircode_by_type(Tuple{typeof(f),Float64})[1][1]  
1 1 - %1 = invoke Main.cos(_2::Float64)::Float64  
    | %2 = invoke Main.sin(%1::Float64)::Float64  
    | return %2
```

- ▶ Data must be stored on the forward-pass of reverse mode, and recovered on the reverse pass
- ▶ Zygote does not have type information
- ▶ `Core.OpaqueClosure` unavailable to Zygote

# Mutation

```
function f(x::Vector{Float64})  
    x .= 5 * x  
    return sin(x[1])  
end
```

- ▶ There *must* be a differentiable function associated to  $f$
- ▶ What is it?
- ▶ Central trick: model state transition<sup>1</sup>

$$\Phi(x) = (5x, \sin(5x_1))$$

- ▶ Undo changes to state on the reverse-pass

---

<sup>1</sup>Pearlmutter and Siskind 2008.

# Supported Language Features

## Supported Language Features

```
f(x::AbstractVector{<:Real}) = sum(5x)
```

Mooncake ✓, Enzyme, ✓, ReverseDiff ✓, Zygote ✓

# Supported Language Features

```
function f(x::AbstractVector{<:Real})  
    for n in eachindex(x)  
        x[n] = exp(x)  
    end  
    return x  
end
```

Mooncake ✓, Enzyme ✓, ReverseDiff ✓, Zygote ✗

# Supported Language Features

```
function f(x::Vector{Float64})  
    for n in eachindex(x)  
        x[n] = exp(x)  
    end  
    return x  
end
```

Mooncake ✓, Enzyme ✓, ReverseDiff ✗, Zygote ✗

# Supported Language Features

- ▶ Complicated Types ✓
- ▶ Incomplete Initialisation ✓
- ▶ Self-referential data structures ✓
- ▶ Dynamic dispatch ✓
- ▶ Value-dependent control flow ✓
- ▶ BLAS (mostly) / LAPACK (somewhat) ✓
- ▶ Structured matrices (Diagonal, Symmetric...) ✓
- ▶ Callable structs / closures ✓



# Limitations

# Basic Limitations

Julia

- ▶ Exception handling (`UpsilonNodes` and `PhiCNodes`)
- ▶ `foreigncalls` always require rules
- ▶ Higher memory usage (in general)
- ▶ Self-referential types

```
struct F
    x::Union{Nothing,F}
end
```

# Current Limitations: Cheap Scalar Operations

- ▶ Vectorisation (SIMD) obstructed
- ▶ Sub-optimal storage strategies (loop invariants, induction variables, constants, etc)
- ▶ See issue 156 for thoughts.

# Conclusion

- ▶ Challenges remain
- ▶ Some important performance and language feature support limitations of RD / Zygote addressed my Mooncake
- ▶ Written in, and operates on, Julia